

B-H for incentive calculation

Contents

Setup params	1
p-value PDF, CDF, invCDF defs	1
Expected value of number of selections	2
Expected total payout	3
Grid search for the good alpha	4
Comparing simulation to experimental stimuli data	10

```
library(ggplot2)
library(dplyr)
library(purrr)
library(furrr)
library(beepR)
library(tidyR)
library(forcats)
library(patchwork)
```

Purpose of this document: Using Benjamini-Hochberg to justify the FP penalty being 19. The simulation will give us the expected total payout under B-H and FP penalty = 19, so we can create a grid of α 's and $P(\text{null})$'s and see which α gives the highest payout.

Executive summary: under B-H, $\alpha = 0.01$ (the lowest tested on the simulation) gives the highest expected total payout. Without any correction strategy, $\alpha = 0.17 \sim 0.2$ maximizes payout.

Setup params

```
alpha <- 0.05      # not using in simulation
K <- 20            # number of data points in a hypothesis/region

# ACHTUNG: mu/sd is c(1.2/3, 1.5/2.5, 1.6/2) in stimuli, picking the average here
mu <- 1.5          # sample mean in a hypothesis/region
sigma <- 2.5

p_null <- 0.5      # the bane of my existence
n_iter <- 5000     # number of iterations in simulation
```

p-value PDF, CDF, invCDF defs

- Definitions of PDF and CDF of the p -value from [Hung_Behavior_1997]
- Random draws from the p -value PDF \mathbf{rp} is sampled through random draws from the uniform $[0, 1]$ quantile space, then looked up through the inverse CDF function (using `uniroot`).
- **Limitation:** These functions are dependent on the μ, σ, K parameters. These parameters are most likely different IRL but we are not using other values yet.

```
# PDF
f_p <- function(x, mu, sigma, K) {
  dnorm(qnorm(1 - x) - sqrt(K) * mu / sigma) / dnorm(qnorm(1 - x))
}

# CDF
F_p <- function(x, mu, sigma, K) {
```

```

  1 - pnorm(qnorm(1 - x) - sqrt(K) * mu / sigma)
}

# inverse CDF of p-value
F_p_inv <- function(q, mu, sigma, K, l = 0, u = 1){
  uniroot(function(p) F_p(p, mu, sigma, K) - q, lower = l, upper = u)$root
}

# random sample from PDF of p-value using its invCDF
rp <- function( mu, sigma){
  q <- runif(1)
  F_p_inv(q, mu, sigma, K)
}

```

Expected value of number of selections

We can get the expected number of selections under B-H with a few simulation iterations and take the average. Given that all the parameters, such as α , are fixed, the uncertainty comes from the sampling of true μ 's from a Binomial and the random sampling of p -value.

- The simulation has `n_iter` of “trials” in our experiment
- For each trial, draw 8 or 12 true μ (0 or μ) from $Bin(n, 1 - P(null))$
- With the μ and pre-specified σ et al., draw 8 or 12 p -values
- Do the B-H and reject hypotheses/regions accordingly
- We get the number of rejections that *should* happen under B-H, for both 8 and 12 regions.

```

(df <-
  expand.grid(
    nregions = c(8, 12),
    iter = 1:n_iter
  ) %>%
  uncount(nregions, .remove = FALSE, .id = "panel") %>%
  group_by(iter, nregions) %>%
  mutate(
    mu = mu * rbinom(nregions, 1, 1 - p_null),
    p_raw = map_dbl(mu, ~rp(.x, sigma)),
    p_bh = p.adjust(p_raw, method = "BH"),
  ) %>%
  pivot_longer(starts_with("p_"), names_to = "method", values_to = "p") %>%
  mutate(
    true = mu == 0, # null hypothesis being true
    reject = p < alpha
  ) %>%
  group_by(iter, nregions, method) %>%
  summarize(
    tp = sum(!true & reject),
    fp = sum(true & reject),
    tn = sum(true & !reject),
    fn = sum(!true & !reject),
    .groups = "drop_last") %>%
  mutate(fdr = ifelse(tp * fp != 0, (fp) / (tp + fp), 0),
    pay = (tp - 19 * fp + tn - fn) * 1)
)

```

A tibble: 20,000 x 9

```
## # Groups:   iter, nregions [10,000]
##   iter nregions method    tp    fp    tn    fn    fdr    pay
##   <int>   <dbl> <chr>   <int> <int> <int> <int> <dbl> <dbl>
## 1     1       8 p_bh     0     0     4     4     0     0
## 2     1       8 p_raw     3     0     4     1     0     6
## 3     1      12 p_bh     4     0     8     0     0    12
## 4     1      12 p_raw     4     0     8     0     0    12
## 5     2       8 p_bh     1     0     5     2     0     4
## 6     2       8 p_raw     2     0     5     1     0     6
## 7     2      12 p_bh     3     0     7     2     0     8
## 8     2      12 p_raw     3     0     7     2     0     8
## 9     3       8 p_bh     4     0     4     0     0     8
## 10    3       8 p_raw     4     0     4     0     0     8
## # ... with 19,990 more rows
```

```
df %>%
  group_by(method) %>%
  summarize(mean(fdr))
```

```
## `summarise()` ungrouping output (override with `.groups` argument)
```

```
## # A tibble: 2 x 2
##   method `mean(fdr)`
##   <chr>         <dbl>
## 1 p_bh         0.0227
## 2 p_raw         0.0523
```

Once we have the distribution of # of rejections B-H says we should make, we can take the average and get the expected value for number of selections.

```
(expected_nselect <-
  df %>%
  group_by(method, nregions) %>%
  summarise(E_nselect = mean(tp + fp), .groups = "drop")
)
```

```
## # A tibble: 4 x 3
##   method nregions E_nselect
##   <chr>     <dbl>     <dbl>
## 1 p_bh         8       3.10
## 2 p_bh        12       4.55
## 3 p_raw         8       3.62
## 4 p_raw        12       5.40
```

Expected total payout

There have been three iterations of expected total payout

1. $E[n_{FP}] = E[n_{reject}] * P(null)$. This is (even more) problematic now that we're using B-H. It can overestimate n_{FP} because if people are ordering hypotheses/regions by p -values, they should be making false discoveries at a lower rate than $P(null)$.
 2. $E[n_{FP}] = E[n_{reject}] * P(true|reject) = E[n_{reject}] * \alpha$. By definition, $P(true|reject) = E\left[\frac{n_{true \wedge reject}}{n_{reject}}\right] = FDR \leq \alpha$, and we say B-H controls FDR at level α . But α here is an upperbound instead of the expected value, see [Benjamini-controlling_1995].
- The problem is, we don't have a good expression for $E[n_{TP}]$. If we write $E[n_{TP}] = E[n_{reject}] * P(\neg true|reject)$, we don't have an expression for $P(\neg true|reject)$ like $FDR \leq \alpha$. If we just use the

joint probability $E[n_{TP}] = nP(\neg true, reject)$, $P(\neg true, reject) = 1 - \beta$, but there is no closed-form for power for B-H; [benjamini_controlling_1995] had to run a simulation. So...

3. Screw it, just use the simulation results. Basically the original simulation.

```
# (expected_payout <- expected_nselect %>%
#   mutate(E_tp = E_nselect * (1 - alpha),
#          E_tn = (nregions - E_nselect) * alpha,
#          E_fp = E_nselect * p_null,
#          E_fn = (nregions - E_nselect) * (1 - p_null),
#          E_payout = E_tp - 19 * E_fp + E_tn - E_fn)
# )

# (expected_payout <- expected_nselect %>%
#   mutate(E_tp = E_nselect * (1 - p_null),
#          E_tn = (nregions - E_nselect) * p_null,
#          E_fp = E_nselect * p_null,
#          E_fn = (nregions - E_nselect) * (1 - p_null),
#          E_payout = E_tp - 19 * E_fp + E_tn - E_fn)
# )
df %>%
  group_by(method, nregions) %>%
  summarize(mean(pay))

## `summarise()` regrouping output by 'method' (override with `.groups` argument)

## # A tibble: 4 x 3
## # Groups:   method [2]
##   method nregions `mean(pay)`
##   <chr>      <dbl>      <dbl>
## 1 p_bh         8         4.14
## 2 p_bh        12         6.37
## 3 p_raw         8         2.93
## 4 p_raw        12         4.22
```

Grid search for the good alpha

Putting the above pieces together, we vary α and $P(null)$ but keep μ , σ and K the same.

Encapsulate simulation function

This is the same code as above

```
finding_payout <- function(alpha, p_null){
  df <-
    expand_grid(
      nregions = c(8, 12),
      iter = 1:n_iter
    ) %>%
    uncount(nregions, .remove = FALSE, .id = "panel") %>%
    group_by(iter, nregions) %>%
    mutate(
      mu = mu * rbinom(nregions, 1, 1 - p_null),
      p_raw = map_dbl(mu, ~rp(.x, sigma)),
      p_bh = p.adjust(p_raw, method = "BH"),
    ) %>%
    pivot_longer(starts_with("p_"), names_to = "method", values_to = "p") %>%
```

```

mutate(
  true = mu == 0, # null hypothesis being true
  reject = p < alpha
) %>%
group_by(iter, nregions, method) %>%
summarize( tp = sum(!true & reject),
           fp = sum(true & reject),
           tn = sum(true & !reject),
           fn = sum(!true & !reject),
           .groups = "drop_last") %>%
mutate(fdr = ifelse(tp * fp != 0, (fp) / (tp + fp), 0),
       pay = (tp - 19 * fp + tn - fn) * 1,
       power = (tp)/(tp + fn)) # this produces NA's; set to 1?

df %>%
group_by(method, nregions) %>%
summarize(E_pay = mean(pay),
          E_fdr = mean(fdr),
          power = mean(power), .groups = "drop")
}

```

setting seed <https://davisvaughan.github.io/furrr/articles/articles/gotchas.html#argument-evaluation>

```

plan(multisession, workers = 5)
options <- furrr_options(seed = 123)

(sim_df <- expand_grid(
  alpha = seq(from = 0.01, to = 0.5, by = 0.01),
  p_null = ppoints(20)
) %>%
split(1:nrow(.)) %>%
future_map_dfr(~cbind(.x, finding_payout(.$alpha, .$p_null), row.names = NULL), .options = options)
)

beep()
saveRDS(sim_df, "sim_n5000_alpha_pnull.rds")

sim_df <- readRDS("sim_n5000_alpha_pnull.rds")

```

Results

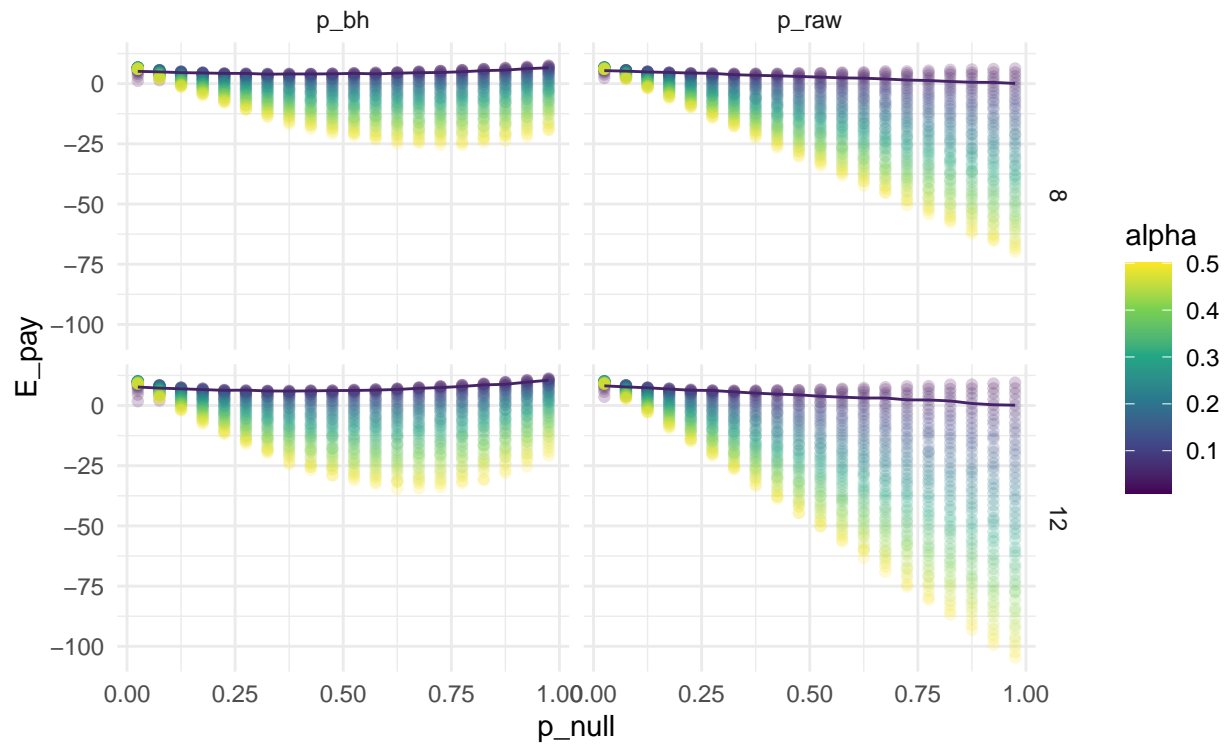
```

sim_df %>%
ggplot(aes(p_null, E_pay, color = alpha)) +
geom_point(alpha = 0.2) +
geom_line(data = sim_df %>% filter(alpha == 0.05)) +
facet_grid(nregions ~ method) +
scale_color_viridis_c() +
theme_minimal() +
labs(title = "Which alpha maximizes payout?", subtitle = "Line drawn at alpha = 0.05")

```

Which alpha maximizes payout?


Line drawn at alpha = 0.05



```
sim_df %>%
  filter(alpha < 0.2 & method != "p_raw") %>%
  ggplot(aes(p_null, E_pay, color = alpha)) +
  geom_point(alpha = 0.2) +
  geom_line(data = sim_df %>% filter(alpha == 0.05 & method != "p_raw")) +
  facet_grid(nregions ~ method) +
  scale_color_viridis_c() +
  theme_minimal() +
  labs(title = "Which alpha maximizes payout?", subtitle = "Line drawn at alpha = 0.05")
```

Which alpha maximizes payout?

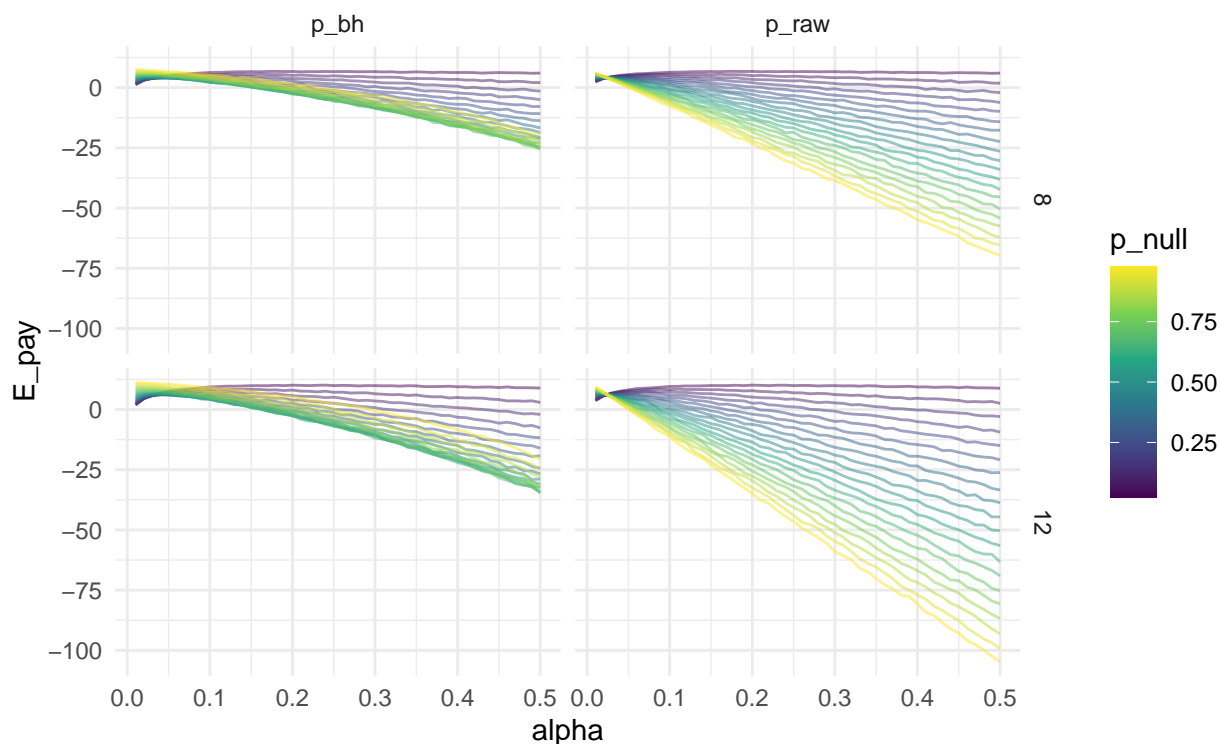
Line drawn at alpha = 0.05



```
sim_df %>%
  ggplot(aes(alpha, E_pay, color = p_null)) +
  # geom_point(alpha = 0.2) +
  geom_line(aes(group = p_null), alpha = 0.5) +
  # geom_line(data = sim_df %>% filter(p_null %in% c(0.025, 0.475, 0.975)), aes(group = p_null)) +
  facet_grid(nregions ~ method) +
  scale_color_viridis_c() +
  theme_minimal() +
  labs(title = "Which alpha maximizes payout?", subtitle = "alpha = 0.01 for B-H, 0.18 for raw?")
```

Which alpha maximizes payout?

alpha = 0.01 for B-H, 0.18 for raw?




```
sim_df %>%
  group_by(nregions, method) %>%
  slice_max(E_pay)
```

```
## # A tibble: 5 x 7
## # Groups:   nregions, method [4]
##   alpha p_null method nregions E_pay   E_fdr   power
##   <dbl> <dbl> <chr>      <dbl> <dbl> <dbl>   <dbl>
## 1  0.01  0.975 p_bh         8  7.54 0.0008   NaN
## 2  0.17  0.025 p_raw         8  6.75 0.00393  0.957
## 3  0.18  0.025 p_raw         8  6.75 0.00441  0.962
## 4  0.01  0.975 p_bh        12 11.3 0.00113   NaN
## 5  0.2   0.025 p_raw        12 10.1 0.00490  0.968
```

```
sim_df %>%
  ggplot(aes(alpha, E_fdr, color = p_null)) +
  geom_point(alpha = 0.2) +
  geom_line(data = sim_df %>% filter(p_null == 0.475)) +
  geom_abline(aes(intercept = 0, slope = 1)) +
  facet_grid(. ~ method) +
  scale_color_viridis_c() +
  theme_minimal() +
  labs(title = "FDR under B-H and raw strategy", subtitle = "Diagonal line is alpha = FDR, colored line
```


FDR under B–H and raw strategy

Diagonal line is $\alpha = \text{FDR}$, colored line at $P(\text{null}) = 0.5$




```
sim_df %>%
  ggplot(aes(alpha, power)) +
  geom_point(aes(color = p_null), alpha = 0.2) +
  geom_vline(aes(xintercept = 0.05)) +
  geom_hline(aes(yintercept = 0.8)) +
  facet_grid(nregions ~ method) +
  scale_color_viridis_c() +
  theme_minimal() +
  labs(title="Power is higher without correction", subtitle = "A bunch of NA's for when P(null) is high")
```

Warning: Removed 2430 rows containing missing values (geom_point).

Power is higher without correction

A bunch of NA's for when P(null) is high



Comparing simulation to experimental stimuli data

```
data.sim.12reg <- readRDS(file = "../data/simulated_data_12regions.rds")
data.sim.8reg <- readRDS(file = "../data/simulated_data_8regions.rds")

(df_stimuli <- bind_rows(
  list("12" = data.sim.12reg, "8" = data.sim.8reg),
  .id = "nregions") %>%
  select(-c(region_idx, effect_size, population, p_h1)) %>%
  mutate(nregions = as.numeric(nregions))
)
```

```
## # A tibble: 700 x 6
##   nregions trial mean    sd    mu data
##   <dbl> <int> <dbl> <dbl> <dbl> <list>
## 1      12     1  0.9    3    0 <dbl [20]>
## 2      12     1  0.9    3    0 <dbl [20]>
## 3      12     1  0.9    3  0.9 <dbl [20]>
## 4      12     1  0.9    3    0 <dbl [20]>
## 5      12     1  0.9    3    0 <dbl [20]>
## 6      12     1  0.9    3    0 <dbl [20]>
## 7      12     1  0.9    3  0.9 <dbl [20]>
## 8      12     1  0.9    3    0 <dbl [20]>
## 9      12     1  0.9    3  0.9 <dbl [20]>
## 10     12     1  0.9    3    0 <dbl [20]>
## # ... with 690 more rows
```

```

stimuli_fn <- function(df_stimuli, alpha){
  df_stimuli %>%
    mutate(p_raw =
      map_dbl(data, ~ t.test(.x, alternative = "greater")$p.value)) %>%
    group_by(nregions, trial) %>%
    mutate(p_bh = p.adjust(p_raw, "BH")) %>%
    select(-data) %>%
    pivot_longer(starts_with("p_"), names_to = "method", values_to = "p") %>%
    mutate(reject = p < alpha,
      true = mu == 0) %>%
    group_by(trial, nregions, method) %>%
    summarize(
      tp = sum(!true & reject),
      fp = sum(true & reject),
      tn = sum(true & !reject),
      fn = sum(!true & !reject),
      .groups = "drop_last") %>%
    mutate(fdr = ifelse(tp * fp != 0, (fp) / (tp + fp), 0),
      pay = (tp - 19 * fp + tn - fn) * 1,
      power = (tp)/(tp + fn)) %>% # this produces NA's; set to 1?
    group_by(nregions, method) %>%
    summarize(E_pay = mean(pay),
      E_fdr = mean(fdr),
      power = mean(power), .groups = "drop")
}


stimuli_alpha_df <- expand.grid(alpha = seq(from = 0.01, to = 0.5, by = 0.01)) %>%
  split(1:nrow(.)) %>%
  map_dfr(~ cbind(.x, stimuli_fn(df_stimuli, .$alpha), row.names = NULL))

(p1 <- stimuli_alpha_df %>%
  ggplot(aes(alpha, E_pay)) +
  geom_point() +
  facet_grid(nregions ~ method) +
  scale_color_viridis_c() +
  theme_minimal() +
  labs(title = "t-test on experiment stimuli", subtitle = "P(null) = 0.5")
)

```


t-test on experiment stimuli

$P(\text{null}) = 0.5$




```
(p2 <- sim_df %>%
  filter(p_null == 0.475) %>%
  ggplot(aes(alpha, E_pay)) +
  # geom_point(alpha = 0.2) +
  geom_point() +
  # geom_line(data = sim_df %>% filter(p_null %in% c(0.025, 0.475, 0.975)), aes(group = p_null)) +
  facet_grid(nregions ~ method) +
  scale_color_viridis_c() +
  theme_minimal() +
  labs(title = "Simulation results")
)
```

Simulation results



p1 + p2

t-test on experiment stimuli
 $P(\text{null}) = 0.5$



Simulation results

